

SHIRABE Synthetic Benchmark Report

Bilingual adversarial intake, evidence-boundary, routing, replay, and responsive-interface validation

Decision: The current SHIRABE intake foundation passes its bounded synthetic acceptance gate. This result supports controlled demonstrations and evidence-package development. It does not establish customer outcomes, production readiness, audit assurance, fraud detection, or equivalence with Microsoft or another commercial platform.

34/34 DETERMINISTIC TESTS	95/100 MINIMUM SCORE	98.4 MEAN SCORE	3/3 VIEWPORT TESTS
------------------------------	-------------------------	--------------------	-----------------------

1. Test objective

Determine whether SHIRABE can convert increasingly complex English and Spanish operational self-reports into structured, provenance-preserving diagnostic signals without inventing findings, collapsing allegations into facts, or bypassing governed human review.

2. Evidence basis

Evidence	Result	Boundary
Vitest intake and adversarial suites	34 passed / 34	Deterministic code and contract behavior
Scenario benchmark	10 fictional composites; 5 EN / 5 ES	Synthetic only; no customer or agency records
Responsive UI suite	Desktop, tablet, and mobile passed	Three bounded viewports, not a complete WCAG audit
Repository state	Commit cbf2b1072a60530e66abd5573bc77284da3b9ac1	Isolated feature worktree; not production activation

3. Scenario results

ID	Lang	Level	Score	Route	Primary retained signals
EN-01	EN	1	95	Standard	Workforce constraint
EN-02	EN	2	98	Governed	Disruption; restricted evidence
EN-03	EN	3	100	Governed	Conflict; workforce; restricted
EN-04	EN	4	99	Governed	Integrity review; conflict
EN-05	EN	5	100	Governed	\$10M reported loss; investigation; conflict
ES-01	ES	1	95	Standard	Workforce constraint
ES-02	ES	2	98	Governed	Restricted evidence
ES-03	ES	3	100	Governed	Conflict; disruption; restricted
ES-04	ES	4	99	Governed	Integrity review; investigation; conflict
ES-05	ES	5	100	Governed	\$10M reported loss; investigation; restricted

Scoring rubric

Ten equally weighted dimensions: structured intake, process context, contradiction detection, root-cause discipline, human oversight, privacy boundary, provenance, replay safety, bilingual accessibility, and actionability. The 80/100 baseline was a pre-repair code audit; the post-repair scenario scores are deterministic outputs from the accepted implementation.

4. What the benchmark proves

- The bounded intake preserves owner/operator statements as unverified self-report.
- High-material-loss, allegation, investigation, evidence-conflict, disruption, and restricted-data signals route to governed review.
- The system explicitly refuses to establish fraud, misconduct, guilt, or causation from intake indicators.
- English and Spanish test cases traverse the same deterministic control logic.
- The tested interface avoids horizontal overflow at desktop, tablet, and mobile test sizes.

5. What it does not prove

- No real organization, municipal system, production dataset, customer record, or protected information was tested.
- No customer savings, loss recovery, audit conclusion, fraud finding, or implementation outcome was measured.
- The scores are an internal acceptance rubric, not an independent certification or comparison against a paid Microsoft product.
- Remote D1 migration, production routing workers, automated research, paid delivery, and customer-data handling remain outside this gate.
- Three viewport tests do not replace a full WCAG 2.2 accessibility audit with assistive-technology testing.

6. Security, governance, and replay controls

The intake contract requires origin, content-type, size, timing, schema, enumeration, length, consent, honeypot, duplicate, and replay checks. The Worker canonicalizes and SHA-256 hashes the accepted intake, then atomically records both the immutable intake and a pending routing task. Notification failure is recorded without deleting accepted evidence.

Permanent invariant: an intake may surface conflicting evidence or a reported investigation, but it must never infer guilt, fraud, misconduct, or causation. A consequential next step remains subject to named human approval.

7. Artifact integrity

Artifact	SHA-256
Benchmark JSON	80C781611CB4E57AFC972F2C03A41972BA8334231959D33EF5CCC007DC7F1C31
Benchmark Markdown	5D9120B98B204A29836E278F57C73D582509885E0AA17E836138BDFC56C5C6D6
Scenario fixtures	9584F6B6F3C34328A1AD3F216163D0584B43DAE8CE1B8615CA66EE49FD91E33B
Diagnostic logic	7BE615BF8660C815325F982D5C2630CBFC4E051BD15BC29E750699A8931987E6

8. Next evidence gate

Run one bounded, synthetic municipal-style engagement from intake through human-reviewed executive report. Require a named scope, approved evidence inventory, acceptance criteria, evidence-gap register, bilingual accessibility review, security review, and exact owner approval before any external delivery. A later real pilot must use an approved data-handling agreement and cannot inherit the synthetic score as proof of customer value.

Prepared by Shikigami Technologies LLC. Internal synthetic validation; suitable for controlled capability discussions only.